Causal Exaggeration

Vincent Bagilet *

Abstract

The credibility revolution in economics has made causal inference methods ubiquitous. At the same time, an increasing amount of evidence has highlighted that this literature strongly favors statistically significant results. I show that these two phenomena interact in a way that significantly worsens the reliability of published estimates: while causal inference methods can successfully reduce bias caused by confounders, the accompanying reduction in statistical power can create another type of bias, leading published estimates to exaggerate true effects sizes. I document this trade-off using realistic fake data simulations and a formal mathematical derivation and discuss potential avenues to address it.

Preliminary introduction

*Columbia University, New York, USA. Email: vincent.bagilet@columbia.edu. A previous version of this paper titled "Unconfounded but Inflated Causal Estimates" was co-authored with Léo Zabrocki-Hallak, I cannot thank him enough for his invaluable contributions to the project. I am very grateful to Hélène Ollivier and Jeffrey Shrader for their guidance and thank Sylvain Chabé-Ferret, Jesse McDevitt-Irwin, David McKenzie, José Luis Montiel Olea, Suresh Naidu, Claire Palandri, Stephan Thies and Roberto Zuniga Valladares for helpful comments, as well as lab members at Columbia and seminars participants at Columbia, IPWSD, the Paris School of Economics and the Toulouse School of Economics.

I Introduction

One of the main challenges of empirical economics is to reduce confounding to identify causal effects. Identifications strategies such as Regression Discontinuity (RD), Instrumental Variable (IV), Difference-in-Differences (DID) and event studies help us achieve this goal. To do so, these strategies only use part of the variation in the data. They exploit the exogenous part of the variation in the treatment or decrease the sample size by only considering observations for which the as-if random assignment assumption is credible. Reducing the variation used can decrease precision and thus statistical power, that is to say the probability of rejecting the null hypothesis of no effect when it is false. In this paper, I argue that this often creates a tension between statistical power and reducing confounding.

When statistical power is low, not only the estimator is imprecise as we well know but statistically significant estimates exaggerate the true effect size (Ioannidis 2008, Gelman and Carlin 2014, Lu et al. 2019, Zwet and Cator 2021). Only estimates at least two standard errors away from zero are statistically significant at the 5% level. In under-powered studies, these estimates make up a selected sub-sample of all estimates, located in the tails of the distribution of all possible estimates. The average of these statistically significant estimates differs from the true effect, located at the center of the distribution if the estimator is unbiased. When power is low, obtaining a statistically significant estimate from an unbiased estimator does not guarantee that it will be close to the true effect. An estimate $\hat{\beta}$ of β might be unbiased in the traditional sense of $\mathbb{E}[\hat{\beta}] = \beta$ but conditionally biased in the sense that $\mathbb{E}[\hat{\beta}|\text{ Significant}] \neq \beta$. Figure 1 illustrates the inflation of significant estimates caused by imprecision.

This consequence of low statistical power could be non-problematic if a large literature had not underlined the existence of a publication bias favoring statistically significant results (Rosenthal 1979, Andrews and Kasy 2019, Abadie 2020, Brodeur et al. 2020, for instance). Published estimates from under-powered studies thus form a biased sample of the distribution of all estimates and can greatly exaggerate true effect sizes. This participates to the current replication crisis affecting various fields such as economics, epidemiology, medicine or psychology (Button et al. 2013, Open Science Collaboration 2015, Camerer et al. 2016, Chang and Li 2022). Even in experimental economics, with a high level of control and an arguable absence of confounders, estimates published in top economic journals that have been replicated were on average inflated by a factor of at least 1.5 (Camerer et al. 2016). Quasi-experimental studies are likely even more exposed to this exaggeration issue as, in current practices, statistical power is not central to observational analyses. Ioannidis et al. (2017) concernedly finds that the median statistical power in a wide range of economic studies is no more than 18% and that nearly 80% of estimates may be exaggerated by a factor of two, despite the widespread use of convincing causal identification strategies and usually large sample

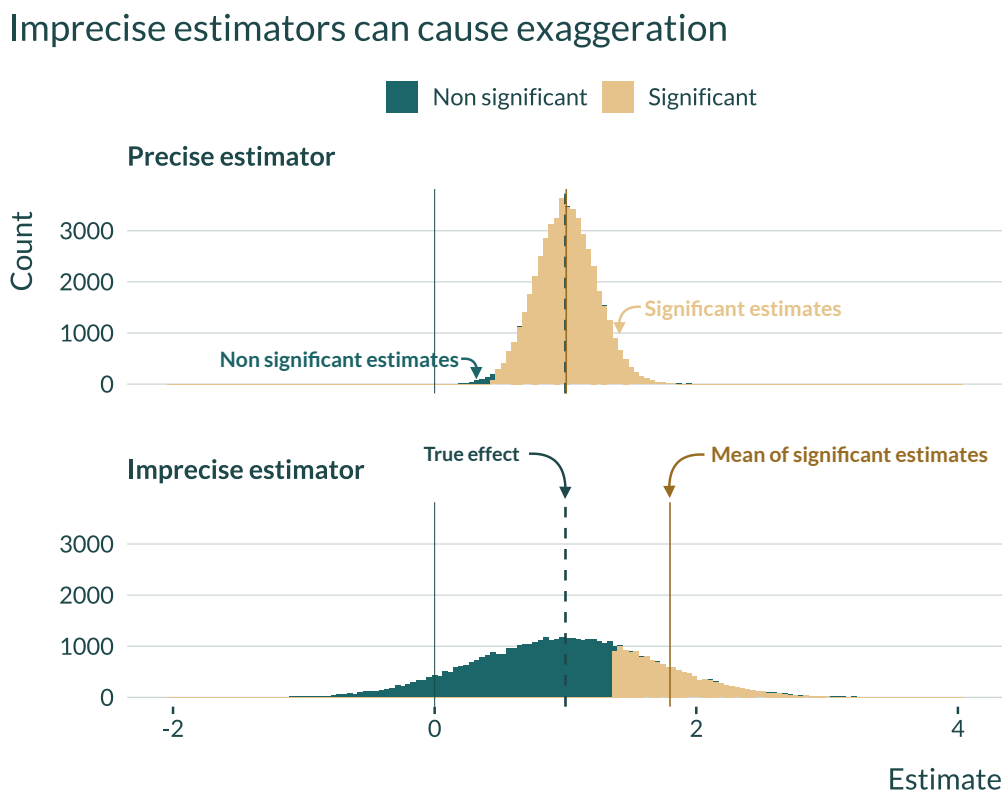

Figure 1 – Significance and distribution of two unbiased estimators with different variances

sizes. The magnitude of exaggerations is thus considerable and in some situations could be on par with that of bias caused by confounders. It is thus crucial to take exaggeration risks into account.

In this paper, I propose an overarching mechanism that may contribute to explaining these inaccuracies in published estimates. I show that design choices in quasi-experimental studies can be seen as a trade-off between avoiding confounding and overestimating true effect sizes due to a resulting loss in power. To limit the threat of confounding, causal inference methods discard variation and therefore reduce statistical power. This can result in inflated significant and published estimates. While causal identifications strategies are essential to describe causal relationships, this paper emphasizes that a perfectly convincing identification does not guaranty an absence of "bias" and that improving identification can actually pull us away from the true effect. In some settings, reaching unbiasedness can actually be impossible as getting rid of bias caused by confounders creates another type of bias.

While the general idea that causal inference methods discard variation to identify effects is shared across strategies, the confounding / exaggeration trade-off is mediated through a distinctive channel for each of them. In RD designs, even when the initial sample size is large, we discard part

of the variation by only considering observations within the bandwidth, decreasing the effective sample size and thus precision. In an IV setting, we only use the subset of the variation in the treatment that is explained by the instrument. In event studies, the variation used to identify an effect sometimes only comes from a limited number of changes in treatment status. When all confounders are measured, matching has a causal interpretation but it still prunes treated units that cannot be matched to untreated ones. Generic controls or fixed effects can increase the variance of the estimate and exaggeration if they absorb more of the variation in the treatment than in the outcome variable.

Since causal identification strategies can be interpreted as ways of controlling for confounders, this last point actually ties all the strategy specific arguments together. We implement Fixed Effects (FEs) based identification strategies such as DiD to control for the invariant, unobserved and arguably endogenous part of the variation in the outcome. In the Control Function (CF) approach of IV, we control for the predicted residuals of the regression of the endogenous variable of interest on the instrument. Fuzzy-RD and propensity score matching can also be thought of as control function approaches, of the forcing variable and propensity score respectively. In addition, excluding observations that are outside the bandwidth or unmatched is equivalent to controlling for observation level fixed effects for these observations. When these methods for controlling for confounders absorb more of the variation in the treatment than in the outcome, they will increase the variance and cause exaggeration. Considering a simple linear homoskedastic model gives the intuition for this trade-off between exaggeration and omitted variable bias for control approaches. Let $y_i = \alpha + \beta x_i + \delta w_i + u_i$, $\forall i \in \{1, .., n\}$, with w potentially unobserved. Under usual assumptions and using the Frisch-Waugh-Lovell theorem we get that σ^2_{OVB} and σ^2_{CTRL}, the variance of the estimators for β when omitting w (short regression) and controlling for it (long regression) respectively are:

$$\sigma^2_{\text{OVB}} = \frac{\sigma^2_{u_{\text{OVB}}}}{n\ \sigma^2_x} = \frac{\sigma^2_{y^{\perp x}}}{n\ \sigma^2_x} \qquad \text{and} \qquad \sigma^2_{\text{CTRL}} = \frac{\sigma^2_{u_{\text{CTRL}}}}{n\ \sigma^2_{x^{\perp w}}} = \frac{\sigma^2_{y^{\perp x,w}}}{n\ \sigma^2_{x^{\perp w}}}$$

and thus,

$$\sigma^2_{\text{OVB}} < \sigma^2_{\text{CTRL}} \quad \Leftrightarrow \quad \frac{\sigma^2_{y^{\perp x}}}{n\ \sigma^2_x} < \frac{\sigma^2_{y^{\perp x,w}}}{n\ \sigma^2_{x^{\perp w}}} \quad \Leftrightarrow \quad \frac{\sigma^2_{x^{\perp w}}}{\sigma^2_x} < \frac{\sigma^2_{y^{\perp x,w}}}{\sigma^2_{y^{\perp x}}}$$

Controlling will increase the variance of the estimator if the fraction of the variance unexplained by w is greater for $y^{\perp x}$ than for x. Put differently, if controlling absorbs more of the variation in x than in the residual part of y ($y^{\perp x}$), it will increase the variance of the estimator. As briefly discussed above, since exaggeration increases with the variance of the estimator, controlling may increase exaggeration. I develop this intuition and a formal proof showing that exaggeration can

be larger when controlling, even when accounting for bias, in section III.

In the remainder of the paper, I first illustrate the existence of this "causal exaggeration" using examples drawn from education, labor, environmental and political economics. The exaggeration of statistically significant estimates can be defined as the ratio of the estimated effect over the true effect, which is never known in a real world setting. To know the true effect and be able to compute this quantity, I have to turn to simulations. Since my fake-data Monte-Carlo simulations have an illustrative purpose only, I intentionally focus on settings in which statistical power can be low. All other simulation assumptions are chosen to make it easier to retrieve the effect of interest. I consider simple linear models with constant treatment effects, i.i.d. observations and homoskedastic errors. All the models are correctly specified and accurately represent the data generating process, except for the omitted variable.

Then, I develop the mathematical intuition behind the confounding / exaggeration trade-off through the tying interpretation of causal identification strategies as control strategies. I also derive a formal proof of the existence of the trade-off for most prevailing causal identification strategies. Specifically, I show that the bias caused by exaggeration can be larger than the one caused by confounders. I also study the drivers of exaggeration and show that it increases as the strength of the instrument decreases, the number of exogenous shocks decreases or when controlling for an omitted variable absorbs more of the variation in the treatment than in the outcome.

Finally, I discuss concrete avenues to address this causal exaggeration when carrying out a non-experimental study. First, I advocate for the use of tools to evaluate the potential magnitude and risk of both confounding and exaggeration issues separately. Sensitivity analyses help with the former while power calculations help with the latter. For instance, the sensitivity analysis tools developed in Cinelli and Hazlett (2020) enable to assess how strong confounders would have to be to change the estimate of the treatment effect beyond a given level we are interested in. Prospective power simulations help identify the design parameters affecting power and exaggeration by approximating the data generating process (Gelman 2020, Black et al. 2021). Retrospective power calculations allow to evaluate whether a study would have enough power to confidently estimate a range of smaller but credible effect sizes (Gelman and Carlin 2014, Stommes et al. 2021). Focusing more specifically on the trade-off and its drivers, I present tools to visualize the variation actually used for identification when using causal identification strategies. The companion website describes in details how such analyses can be implemented. Finally, I briefly discuss potential solutions to mitigate this trade-off.

This paper contributes to three strands of the literature. First, the idea that causal identification estimators, while unbiased, may be imprecise is not new; this is part of the well-known bias / variance trade-off (Imbens and Kalyanaraman 2012, Deaton and Cartwright 2018, Hernán and Robins 2020, Ravallion 2020). I approach this literature from a different angle: through the prism

of statistical power and publication bias. Not only the limited precision resulting from the use of causal identification methods could make it difficult to draw clear conclusions regarding the exact magnitude of the effect but I argue that it might also inherently lead to inflated published effect sizes, creating another "bias".

Second, recent studies discussing the exaggeration of statistically significant estimates due to low power focused on specific causal identification methods separately and usually do not investigate the determinants of this exaggeration (Schell et al. 2018, Black et al. 2021, Stommes et al. 2021, Young 2021). In a companion paper, I consider wide range of causal and non-causal identification strategies and highlight tangible design parameters that can cause exaggeration (**?**). In the present paper, I propose an overarching mechanism, inherent to causal identification strategies as a whole, and that can explain these issues: although each strategy does so through different means, in essence they discard part of the variation, thereby causing exaggeration. This connection could be exacerbated by the fact that, as noted by Brodeur et al. (2020), publication bias is more prevalent for some methods such as the IV.

Third, this study contributes to the literature on reproducibility in economics (Camerer et al. 2016, Ioannidis et al. 2017, Christensen and Miguel 2018, Kasy 2021). The trade-off presented in this paper may be an additional explanation for observing replication failures in empirical economics, despite the widespread use of convincing causal identification methods.

References

Abadie, A. (2020), 'Statistical Nonsignificance in Empirical Economics', American Economic Review: Insights **2**(2), 193–208. 2

Amrhein, V., Trafimow, D. and Greenland, S. (2019), 'Inferential Statistics as Descriptive Statistics: There Is No Replication Crisis if We Don't Expect Replication', The American Statistician **73**(sup1), 262–270.

Andrews, I. and Kasy, M. (2019), 'Identification of and Correction for Publication Bias', American Economic Review **109**(8), 2766–2794. 2

Angrist, J. D. and Pischke, J.-S. (2009), Mostly Harmless Econometrics: An Empiricist's Companion, 1 edition edn, Princeton University Press, Princeton.

Angrist, J. D. and Pischke, J.-S. (2010), 'The Credibility Revolution in Empirical Economics: How Better Research Design Is Taking the Con out of Econometrics', Journal of Economic Perspectives **24**(2), 3–30.

Angrist, J. D. and Pischke, J.-S. (2014), Mastering 'Metrics: The Path from Cause to Effect, Princeton University Press.

Athey, S. and Imbens, G. (2016), 'The Econometrics of Randomized Experiments', arXiv:1607.00698 [econ, stat] .

Black, B. S., Hollingsworth, A., Nunes, L. and Simon, K. I. (2021), Simulated Power Analyses for Observational Studies: An Application to the Affordable Care Act Medicaid Expansion, SSRN Scholarly Paper ID 3368187, Social Science Research Network, Rochester, NY. 5, 6

Brodeur, A., Cook, N. and Heyes, A. (2020), 'Methods Matter: P-Hacking and Publication Bias in Causal Analysis in Economics', American Economic Review **110**(11), 3634–3660. 2, 6

Button, K. S., Ioannidis, J. P. A., Mokrysz, C., Nosek, B. A., Flint, J., Robinson, E. S. J. and Munafò, M. R. (2013), 'Power failure: Why small sample size undermines the reliability of neuroscience', Nature Reviews Neuroscience **14**(5), 365–376. 2

Camerer, C. F., Dreber, A., Forsell, E., Ho, T.-H., Huber, J., Johannesson, M., Kirchler, M., Almenberg, J., Altmejd, A., Chan, T., Heikensten, E., Holzmeister, F., Imai, T., Isaksson, S., Nave, G., Pfeiffer, T., Razen, M. and Wu, H. (2016), 'Evaluating replicability of laboratory experiments in economics', Science **351**(6280), 1433–1436. 2, 6

Card, D. (1993), Using Geographic Variation in College Proximity to Estimate the Return to Schooling, Working Paper 4483, National Bureau of Economic Research.

Chang, A. C. and Li, P. (2022), 'Is Economics Research Replicable? Sixty Published Papers From Thirteen Journals Say "Often Not"', Critical Finance Review **11**. 2

Christensen, G. and Miguel, E. (2018), 'Transparency, Reproducibility, and the Credibility of Economics Research', Journal of Economic Literature **56**(3), 920–980. 6

Cinelli, C. and Hazlett, C. (2020), 'Making sense of sensitivity: Extending omitted variable bias', Journal of the Royal Statistical Society: Series B (Statistical Methodology) **82**(1), 39–67. 5

Cunningham, S. (2021), Causal Inference: The Mixtape, Yale University Press.

Deaton, A. and Cartwright, N. (2018), 'Understanding and misunderstanding randomized controlled trials', Social Science & Medicine **210**, 2–21. 5

Duflo, E., Glennerster, R. and Kremer, M. (2007), Using Randomization in Development Economics Research: A Toolkit, in T. P. Schultz and J. A. Strauss, eds, 'Handbook of Development Economics', Vol. 4, Elsevier, pp. 3895–3962.

Gelman, A. (2020), Regression and Other Stories, Cambridge University Press, Cambridge New York, NY Port Melbourne, VIC New Delhi Singapore. 5

Gelman, A. and Carlin, J. (2014), 'Beyond Power Calculations: Assessing Type S (Sign) and Type M (Magnitude) Errors', Perspectives on Psychological Science **9**(6), 641–651. 2, 5

Hernán, M. A. (2021), 'Causal analyses of existing databases: No power calculations required', Journal of Clinical Epidemiology p. S0895435621002730.

Hernán, M. A. and Robins, J. M. (2020), Causal Inference: What If, boca raton: chapman & hall/crc edn. 5

Hill, J. L. (2011), 'Bayesian Nonparametric Modeling for Causal Inference', Journal of Computational and Graphical Statistics **20**(1), 217–240.

Huntington-Klein, N. (2021), The Effect: An Introduction to Research Design and Causality, first edn, Chapman and Hall/CRC, Boca Raton.

Imbens, G. and Kalyanaraman, K. (2012), 'Optimal Bandwidth Choice for the Regression Discontinuity Estimator', The Review of Economic Studies **79**(3), 933–959. 5

Imbens, G. W. and Rubin, D. B. (2015), Causal Inference for Statistics, Social, and Biomedical Sciences: An Introduction, Cambridge University Press, Cambridge.

Ioannidis, J. P. A. (2008), 'Why Most Discovered True Associations Are Inflated', Epidemiology **19**(5), 640–648. 2

Ioannidis, J. P. A., Stanley, T. D. and Doucouliagos, H. (2017), 'The Power of Bias in Economics Research', The Economic Journal **127**(605), F236–F265. 2, 6

Kasy, M. (2021), 'Of Forking Paths and Tied Hands: Selective Publication of Findings, and What Economists Should Do about It', Journal of Economic Perspectives **35**(3), 175–192. 6

Linden, A. (2019), 'RETRODESIGN: Stata module to compute type-S (Sign) and type-M (Magnitude) errors', Boston College Department of Economics.

Lu, J., Qiu, Y. and Deng, A. (2019), 'A note on Type S/M errors in hypothesis testing', British Journal of Mathematical and Statistical Psychology **72**(1), 1–17. 2

McConnell, B. and Vera-Hernandez, M. (2015), Going beyond simple sample size calculations: A practitioner's guide, Technical report, Institute for Fiscal Studies.

McShane, B. B., Gal, D., Gelman, A., Robert, C. and Tackett, J. L. (2019), 'Abandon Statistical Significance', The American Statistician **73**(sup1), 235–245.

Open Science Collaboration (2015), 'Estimating the reproducibility of psychological science', Science **349**(6251), aac4716. 2

Oster, E. (2019), 'Unobservable Selection and Coefficient Stability: Theory and Evidence', Journal of Business & Economic Statistics **37**(2), 187–204.

Ravallion, M. (2020), Should the Randomistas (Continue to) Rule?, Working Paper 27554, National Bureau of Economic Research. 5

Romer, D. (2020), 'In Praise of Confidence Intervals', AEA Papers and Proceedings **110**, 55–60.

Rosenbaum, P. R. (2002), Observational Studies, Springer Series in Statistics, Springer New York, New York, NY.

Rosenthal, R. (1979), 'The file drawer problem and tolerance for null results', Psychological Bulletin **86**(3), 638–641. 2

Schell, T. L., Griffin, B. A. and Morral, A. R. (2018), Evaluating Methods to Estimate the Effect of State Laws on Firearm Deaths: A Simulation Study, Technical report, RAND Corporation. 6

Shadish, W. R., Cook, T. D. and Campbell, D. T. (2002), Experimental and Quasi-experimental Designs for Generalized Causal Inference, Houghton Mifflin.

Stommes, D., Aronow, P. M. and Sävje, F. (2021), 'On the reliability of published findings using the regression discontinuity design in political science', arXiv:2109.14526 [stat] . 5, 6

Timm, A., Gelman, A. and Carlin, J. (2019), 'Retrodesign: Tools for Type S (Sign) and Type M (Magnitude) Errors'.

Young, A. (2021), 'Leverage, Heteroskedasticity and Instrumental Variables in Practical Application', p. 43. 6

Zwet, E. and Gelman, A. (2021), 'A Proposal for Informative Default Priors Scaled by the Standard Error of Estimates', The American Statistician pp. 1–9.

Zwet, E., Schwab, S. and Senn, S. (2021), 'The statistical properties of RCTs and a proposal for shrinkage', Statistics in Medicine **40**(27), 6107–6117.

Zwet, E. W. and Cator, E. A. (2021), 'The significance filter, the winner's curse and the need to shrink', Statistica Neerlandica **75**(4), 437–452. 2