Causal Exaggeration

Vincent Bagilet *

Abstract

To avoid confounding, causal identification strategies focus on a subset of the variation in the data, the plausibly exogenous part. In this paper, I argue that it can reduce statistical power and lead published estimates to exaggerate true effects sizes. Using realistic fake data simulations and a mathematical derivation, I show that for the main causal strategies, a perfectly convincing identification does not guaranty an absence of "bias" and that improving identification can actually pull us away from the true effect. I then discuss potential avenues to address this issue.

*Columbia University, New York, USA. Email: vincent.bagilet@columbia.edu. A previous version of this project was co-authored with Léo Zabrocki-Hallak, I cannot thank him enough for his invaluable and far-reaching contributions to the project. I am very grateful to Hélène Ollivier and Jeffrey Shrader for their guidance and thank Sylvain Chabé-Ferret, Jesse McDevitt-Irwin, David McKenzie, José Luis Montiel Olea, Suresh Naidu, Claire Palandri, Stephan Thies and Roberto Zuniga Valladares for helpful comments, as well as lab members at Columbia and seminars participants at Columbia, IPWSD, the Paris School of Economics and the Toulouse School of Economics.

I Introduction

One of the main challenges of empirical economics is to reduce confounding to identify causal effects. Identifications strategies such as Regression Discontinuity (RD), Instrumental Variable (IV), Difference-in-Differences (DID) and event studies help achieve this goal. To do so, these strategies only use part of the variation in the data. They exploit the exogenous part of the variation in the treatment or decrease the sample size by only considering observations for which the as-if random assignment assumption is credible. Reducing the variation used can decrease precision and thus statistical power, that is to say the probability of rejecting the null hypothesis of no effect when it is false. In this paper, I argue that it often creates a tension between statistical power and reducing confounding.

When statistical power is low, not only the estimator is imprecise as we well know but statistically significant estimates exaggerate the true effect size (Ioannidis 2008, Gelman and Carlin 2014, Lu et al. 2019, Zwet and Cator 2021). Only estimates at least two standard errors away from zero are statistically significant at the 5% level. In under-powered studies, these estimates make up a selected sub-sample of all estimates, located in the tails of the distribution of all possible estimates. The average of these statistically significant estimates differs from the true effect, located at the center of the distribution if the estimator is unbiased. When power is low, obtaining a statistically significant estimate from an unbiased estimator does not guarantee that it will be close to the true effect. An estimate $\hat{\beta}$ of β might be unbiased in the traditional sense but conditionally biased in the sense that $\mathbb{E}[\hat{\beta}] = \beta$ but $\mathbb{E}[\hat{\beta}|\text{ Significant}] \neq \beta$. Figure 1 illustrates the inflation of significant estimates caused by imprecision.

This consequence of low statistical power could be non-problematic if a large literature had not underlined the existence of a publication bias favoring statistically significant results (Rosenthal 1979, Andrews and Kasy 2019, Abadie 2020, Brodeur et al. 2020, for instance). Published estimates from under-powered studies thus form a biased sample of the distribution of all estimates and can greatly exaggerate true effect sizes. This participates to the current replication crisis affecting various fields such as economics, epidemiology, medicine or psychology (Button et al. 2013, Open Science Collaboration 2015, Camerer et al. 2016, Chang and Li 2022). Even in experimental economics, with a high level of control and an arguable absence of confounders, estimates published in top economic journals that have been replicated were on average inflated by a factor of at least 1.5 (Camerer et al. 2016). Quasi-experimental studies are likely even more exposed to this exaggeration issue as, in current practices, statistical power is not central to observational analyses. Ioannidis et al. (2017) concernedly finds that the median statistical power in a wide range of economic studies is no more than 18% and that nearly 80% of estimates may be exaggerated by a factor of two, despite the widespread use of convincing causal identification strategies and usually large sample

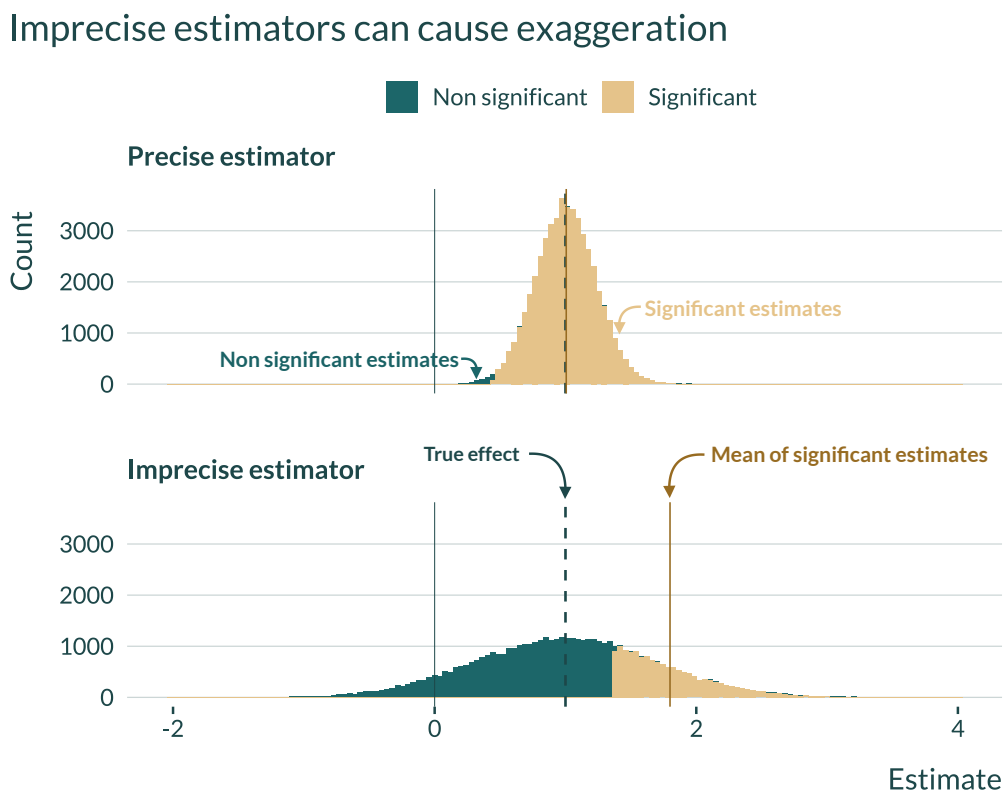

Figure 1 – Significance and distribution of two unbiased estimators with different variances

sizes. The magnitude of exaggerations is thus considerable and in some situations could be on par with that of bias caused by confounders. It is thus crucial to take exaggeration risks into account.

In this paper, I propose an overarching mechanism that may contribute to explaining these inaccuracies in published estimates. I show that design choices in quasi-experimental studies can be seen as a trade-off between avoiding confounding and overestimating true effect sizes due to a resulting loss in power. To limit the threat of confounding, causal inference methods discard variation and therefore reduce statistical power. This can result in inflated significant and published estimates. While causal identifications strategies are essential to describe causal relationships, this paper emphasizes that a perfectly convincing identification does not guaranty an absence of "bias" and that improving identification can actually pull us away from the true effect. In some settings, reaching unbiasedness can actually be impossible as getting rid of bias caused by confounders creates another type of bias.

While the general idea that causal inference methods discard variation to identify effects is shared across strategies, the confounding / exaggeration trade-off is mediated through a distinctive channel for each of them. In RD designs, even when the initial sample size is large, we discard part

of the variation by only considering observations within the bandwidth, decreasing the effective sample size and thus precision. In an IV setting, we only use the subset of the variation in the treatment that is explained by the instrument. In event studies, the variation used to identify an effect sometimes only comes from a limited number of treated observations. Assuming that all confounders are measured, matching prunes treated units that cannot be matched to untreated ones. Generic controls or fixed effects can increase the variance of the estimate and exaggeration if they absorb more of the variation in the treatment than in the outcome variable.

Since causal identification strategies can be interpreted as ways of controlling for confounders, this last point actually ties all the strategy specific arguments together. We implement Fixed Effects (FEs) based identification strategies such as DiD to control for the invariant, unobserved and arguably endogenous part of the variation in the outcome. In the Control Function (CF) approach of IV, we control for the predicted residuals of the regression of the endogenous variable of interest on the instrument. Fuzzy-RD and propensity score matching can also be thought of as control function approaches, of the forcing variable and propensity score respectively. In addition, excluding observations that are outside the bandwidth or unmatched is equivalent to controlling for observation level fixed effects for these observations. Considering a simple linear homoskedastic model gives the intuition for the trade-off between exaggeration and omitted variable bias for control approaches. Let $y_i = \alpha + \beta x_i + \delta w_i + u_i$, $\forall i \in \{1, .., n\}$, with w potentially unobserved. Under usual assumptions and using the Frisch-Waugh-Lovell theorem we get that σ^2_{OVB} and σ^2_{CTRL}, the variance of the estimators for β when omitting w (short regression) and controlling for it (long regression) respectively are:

$$\sigma^2_{\text{OVB}} = \frac{\sigma^2_{u_{\text{OVB}}}}{n\ \sigma^2_x} = \frac{\sigma^2_{y^{\perp x}}}{n\ \sigma^2_x} \qquad \text{and} \qquad \sigma^2_{\text{CTRL}} = \frac{\sigma^2_{u_{\text{CTRL}}}}{n\ \sigma^2_x} = \frac{\sigma^2_{y^{\perp x,w}}}{n\ \sigma^2_{x^{\perp w}}}$$

and thus,

$$\sigma^2_{\text{OVB}} < \sigma^2_{\text{CTRL}} \quad \Leftrightarrow \quad \frac{\sigma^2_{y^{\perp x}}}{n\ \sigma^2_x} < \frac{\sigma^2_{y^{\perp x,w}}}{n\ \sigma^2_{x^{\perp w}}} \quad \Leftrightarrow \quad \frac{\sigma^2_{x^{\perp w}}}{\sigma^2_x} < \frac{\sigma^2_{y^{\perp x,w}}}{\sigma^2_{y^{\perp x}}}$$

Controlling will increase the variance of the estimator if the fraction of the variance unexplained by w is greater for $y^{\perp x}$ than for x. Put differently, if controlling absorbs more of the variation in x than in the residual part of y ($y^{\perp x}$), it will increase the variance of the estimator. As briefly discussed above, since exaggeration increases with the variance of the estimator, controlling may increase exaggeration. I develop this intuition and a formal proof showing that exaggeration can be larger when controlling, even when accounting for bias, in section III.

In the remainder of the paper, I first illustrate the existence of this "causal exaggeration" using examples drawn from education, labor, environmental and political economics. The inflation

of statistically significant estimates can be defined as the ratio of the estimated effect over the true effect, which is never known in a real world setting. I therefore have to turn to fake-data simulations. Since these Monte-Carlo simulations have an illustrative purpose only, I intentionally focus on settings in which statistical power can be low. All other simulation assumptions are chosen to make it easier to retrieve the effect of interest. I consider simple linear models with constant treatment effects, iid observations and homoskedastic errors. All the models are correctly specified and accurately represent the data generating process, except for the omitted variable.

Then, I develop the mathematical intuition behind the confounding / exaggeration trade-off through the tying interpretation of causal identification strategies as control strategies. I also derive a formal proof of the existence of the trade-off for most prevailing causal identification strategies. Specifically, I show that the exaggeration bias caused can be larger than the one caused by confounders. I also study the drivers of exaggeration and show that it increases as the strength of the instrument decreases, the number of exogenous shocks decreases or when controlling for an omitted variable absorbs more of the variation in the treatment than in the outcome.

Finally, I discuss avenues to address this causal exaggeration. First, I advocate for the use of tools to evaluate the potential magnitude and risk of both confounding and exaggeration issues separately. Sensitivity analyses help with the former while power calculations help with the latter. For instance, the sensitivity analysis tools developed in Cinelli and Hazlett (2020) enable to assess how strong confounders would have to be to change the estimate of the treatment effect beyond a given level we are interested in. Prospective power simulations help identify the design parameters affecting power and exaggeration by approximating the data generating process (Gelman 2020, Black et al. 2021). Retrospective power calculations allow to evaluate whether a study would have enough power to confidently estimate a range of smaller but credible effect sizes (Gelman and Carlin 2014, Stommes et al. 2021). The companion website describes in details how such solutions can be implemented. Focusing more specifically on the trade-off and its drivers, I present tools to visualize the variation actually used for identification when using causal identification strategies. Finally, I briefly discuss potential solutions to mitigate this trade-off.

This paper contributes to three strands of the literature. First, the idea that causal identification estimators, while unbiased, may be imprecise is not new; this is part of the well-known bias / variance trade-off (Imbens and Kalyanaraman 2012, Deaton and Cartwright 2018, Hernán and Robins 2020, Ravallion 2020). In under-powered studies, resulting estimates have large confidence intervals, suggesting that a wide range of effects are consistent with the data. I approach this literature from a different angle: through the prism of statistical power and publication bias. Not only the limited precision resulting from the use of causal identification methods could make it difficult to draw clear conclusions regarding the exact magnitude of the effect but I argue that it might also inherently lead to inflated published effect sizes, creating another "bias".

Second, recent studies discussing the inflation of statistically significant estimates due to low power focused on specific causal identification methods separately and usually do not investigate the determinants of this exaggeration (Schell et al. 2018, Black et al. 2021, Stommes et al. 2021, Young 2021). In this paper, I propose an overarching mechanism, inherent to causal identification strategies as a whole, and that can explain these issues: although each strategy does so through different means, in essence they discard part of the variation, thereby causing exaggeration. This connection could be exacerbated by the fact that, as noted by Brodeur et al. (2020), publication bias is more prevalent for some methods such as the IV.

Third, this study contributes to the literature on reproducibility in economics (Camerer et al. 2016, Ioannidis et al. 2017, Christensen and Miguel 2018, Kasy 2021). The trade-off presented in this paper may be an additional explanation for observing replication failures in empirical economics, despite the widespread use of convincing causal identification methods.

II Simulations

III Mathematical derivation

IV Practical Recommendations

We have seen that using causal methods induces a trade-off between avoiding confoundings and exaggerating true effects if the study is under-powered. How can we, as applied researchers using observational data, arbitrate it? Since key pieces of information such as the true effect and the effect of omitted variables are inherently unknown, we cannot directly compute the biases caused by confounders and exaggeration. In this section, I examine how we however can get a sense of threats from both sides of the trade-off and visualize its main driver, the variation used for identification. I then discuss how changing our attitude towards statistical significance and replicating studies could limit the exaggeration issue.

IV. 1 Gauging omitted variable bias

On one side of the trade-off lies the bias caused by confounders. Although it is in essence impossible to measure, tools such as sensitivity analyses are available to gauge its magnitude (Rosenbaum 2002, Oster 2019, Cinelli and Hazlett 2020). For instance, the method developed in Cinelli and Hazlett (2020) enables to assess how strong confounders would have to be to change the estimate of the treatment effect beyond a given level we are interested in. It offers bounds for the strength of the association between the treatment and potential omitted variables by weighting

it against the measured association between the treatment and observed covariates. A typical conclusion from such an analysis would be: "omitted variables would have to explain as much residual variance of the outcome and the treatment as the observed covariate x to bring down the estimate to a value of β_l". In addition, the authors implement graphical tools to facilitate this comparison. I suggest to use such quantitative bias analyses to evaluate the restrictiveness of the causal approach required to limit the threat of unobserved confounding to acceptable levels.

IV. 2 Evaluating risks of exaggeration

On the other side of the trade-off lies the exaggeration emerging when statistical power is low. As OVB, exaggeration and statistical power are in essence impossible to measure as their computation depends on the true effect which is always unknown. Yet, power calculations can help assess them by making hypothesizes on the magnitude of the true effect. In randomized controlled trials, such computations are not only an established practice but a requirement (Duflo et al. 2007, McConnell and Vera-Hernandez 2015, Athey and Imbens 2016). They are, however, rarely reported in non-experimental studies. This can be understandable if power calculations are perceived only as a way to measure a design's ability to detect an effect when there is actually one. Yet, taking publication bias and the threat of exaggeration into account highlights the necessity of running power calculations in non-experimental studies as well. A low power or a relatively large variance not only makes it more difficult to detect an effect or to draw clear conclusions about its magnitude when detected but it can also create a bias. Considering the attention given to bias avoidance in the economics literature, I advocate to make power central to non-experimental analyses to limit bias caused by exaggeration. Currently, in causal inference textbooks, very few pages are devoted to statistical power in non-experimental studies (Angrist and Pischke 2009; 2014, Imbens and Rubin 2015, Cunningham 2021). To the best of my knowledge, only two textbooks discuss the matter in depth (Shadish et al. 2002, Huntington-Klein 2021). Results from power and exaggeration calculations would not only be highly informative but could also be reported very concisely in the robustness section of articles.

IV. 2.1 Prospective power calculations

To evaluate the statistical power of a study, the risk of exaggeration and identifying the factors driving it, one can first simulate the design of the study (Hill 2011, Gelman 2020, Black et al. 2021). Simulating a data generating process from scratch requires thinking about the distribution of the variables, about their relationships and can also help underline the variation used for identification. I implemented such fake-data simulations in Section II and provide `R` code to be used as examples in the replication material. In situations where the relationships among covariates are too complex

to emulate, one can also start from an existing dataset and add a known treatment effect. I implemented real data simulations in a companion paper and describe their implementation in its replication material (Bagilet 2023).

IV. 2.2 Retrospective power calculations

Running post-analysis power calculations can also help getting a sense of the statistical power associated with a research design. Such *retrospective* calculations allow to evaluate whether the design of the study would produce accurate and uninflated statistically significant estimates if the true effect was in fact smaller than the observed estimate (Gelman and Carlin 2014, Ioannidis et al. 2017, Stommes et al. 2021).

I illustrate how a retrospective analysis works by taking the example of Card (1993) on the relationship between human capital and income. He finds that an additional year of education, instrumented by the distance of growing near a four-year college, causes a 13.2% average increase in wage. The associated standard error is 5.5%. Is there a risk of exaggeration with this design? Since, as noted by the author himself, the estimate is very imprecise we could expect so. Imagine that prior evidence suggests that the true effect is likely close to a 10% increase in wage. We can then compute the statistical power of the study by drawing many estimates from a normal distribution centered around the hypothesized true effect of 10% and with a standard deviation equal to the 5.5% standard error obtained in Card (1993). Concretely, this proceeds as if we were able to replicate the study many times under the assumption that the true effect was equal to the hypothesized one [1]. Statistically significant estimates (at the 5% level) would on average be roughly equal to 15%, therefore overestimating the true effect by a factor of 1.5. Statistical power, the proportion of estimates that are significant, would be only 44%. Conditional on a 10% true effect size being a reasonable assumption, this study would be under-powered and the risk of exaggeration large.

The usefulness of any retrospective power analysis lies on the assumption regarding the true effect size. Informed guesses about the range of plausible effect sizes can be based on results from meta-analyses or previous studies with a convincing design (*e.g.*, a large randomized controlled trial) [2]. When such information is not available, power calculations can be ran for a range of smaller but credible effect sizes.

1. Timm et al. (2019), Linden (2019) offer `R` and `Stata` packages that enable to easily run these calculations through a short command: `retrodesign(10, 5.5)`.

2. Note that when such meta-analyses are available, one can use a Bayesian procedure to shrink statistically significant estimates based on the corpus of estimates from prior studies Zwet and Gelman (2021), Zwet and Cator (2021), Zwet et al. (2021) .

IV. 3 Driver of the trade-off

We have seen that the variation used for identification was the main driver of the confounding / exaggeration trade-off. This variation determines the variance of the estimator and thus the amount of exaggeration. In this section, I first discuss the impact of estimator variance and how exaggeration leads to reinterpret the well-known bias/variance trade-off. I then propose to use a tool that enables to visualize where the variation used for identification comes from.

IV. 3.1 Navigating the bias/variance trade-off

In non-experimental studies, estimator variance is often important to the extent that a large variance may lead to a failure to reject the null of zero effect when it is actually incorrect. It makes variance paramount as long as a statistically significant estimate is not obtained. Yet, exaggeration underlines that variance matters, even after that a significant estimate has been obtained.

Obtaining a statistically significant estimate from an imprecise estimator should not necessarily be interpreted as a sign of "success" in getting significance despite the estimator being imprecise. It could however be a warning that this estimate comes from the tails of the distribution and thus inaccurately represents the true effect. Conditional on having obtained a statistically significant estimate, a limited precision can hide a bias: exaggeration. This invites to revisit the well-known bias/variance trade-off: a larger variance can actually lead to a larger (conditional) bias. The bias/variance trade-off can in fact be a bias/bias trade-off. This paper thus invites to pay more attention to the implications of our design choices on the variance of our estimators.

IV. 3.2 Visualizing the driver of the trade-off

Since the variation used for identification is the key driver of the trade-off, visualizing it can be helpful to assess the proportion of the variation (and observations) that actually contributes to the estimation. To do so, I both leverage the interpretation of causal inference methods as control approaches extensively discussed above and build on an approach developed in Aronow and Samii (2016) for another purpose altogether: evaluating external validity of standard regression. Aronow and Samii (2016) essentially interprets the coefficient of a treatment of interest of a simple linear non-causal regression as a weighted average of individual treatment effects. The weight w_i of individual i is simply the squared difference of its treatment T_i and the value of this treatment as predicted by the other covariates X. Thus $w_i = (T_i - \mathbb{E}[T_i|X_i])^2$. The authors use this to show that if treatment effects are heterogenous, the weighting may lead to disproportionally represent some observations when computing the average effect of the treatment. Some observations, whose treatment status is very well explained by covariates do not actually contribute to the estimation of the treatment effect.

The parallel with our settings is now obvious
[3]

We can also reuse their concept of effective sample

IV. 4 Attitude Towards Statistical Significance and Replication

Exaggeration only arises in the presence of publication bias. As shown in our simulations, if estimates were not filtered by their statistical significance, even under-powered studies would on average recover the true effect, as long as the estimator is unbiased.

First, it might be helpful to discuss how a statistically significant result can be obtained.

In low power situations, statistically significant estimates can be obtained either by chance or as an outcome of the garden of forking paths (Simmons et al. 2011, Gelman and Loken 2013, Kasy 2021). Forking paths arise in various decisions during a non-experimental study, such as those regarding data preparation, the inclusion of control variables, and whether editors decide to publish a paper or not. Due to the structural flaw that favors significance, the path chosen in such situations may be more likely to lead to a statistically significant result.

Since exaggeration only arises in thee getting rid of publication bias would

General changes in scientific practices could also limit the inflation of statistically significant estimates in under-powered studies

The publication bias arising from dichotomizing evidence according to p-values has long been criticized in many disciplines but has seen a revival with the recent replication crises in psychology, medicine and social sciences. Many researchers advocate abandoning statistical significance as a measure of a study's quality (McShane et al. 2019). This would essentially eliminate the trade-off described in this paper.

To be effective, this change in attitude towards statistical significance should be paired with an effort to replicate studies (Christensen and Miguel 2018). Replications, even of low powered studies, would eventually enable to build the distribution of the causal estimand of interest. Meta-analyzes would then reduce the uncertainty around the true value of the causal estimand by pooling estimates (Hernán 2021).

Finally, the inflation of statistically significant estimates can be limited by considering confidence intervals as compatibility intervals (Shadish et al. 2002, Amrhein et al. 2019, Romer 2020). The width of these intervals gives a range of effect sizes compatible with the data. Confidence intervals will be wide in under-powered studies signaling that point estimates should not be taken

3. Note that in most of our settings, treatment effect is homogenous and thus the weighted average of individual treatment effects will not depend on these weights. However, the variance of the estimator will.

at face value, even if statistically significant.

V Conclusion

In the presence of publication bias, estimates published from underpowered studies exaggerate true effect sizes, even when the estimators are "unbiased" in the usual sense of $\mathbb{E}[\hat{\beta}] = \beta$ (Ioannidis 2008, Gelman and Carlin 2014, Lu et al. 2019, Zwet and Cator 2021). Ioannidis et al. (2017) showed that the economic literature suffers from an extensive lack of power and many other papers showed the existence of a publication bias in this literature (Rosenthal 1979, Andrews and Kasy 2019, Abadie 2020, Brodeur et al. 2020, for instance). It is therefore not surprising that many estimates published in economics have been shown to be considerably exaggerated (Camerer et al. 2016, Ioannidis et al. 2017), despite the extensive use of convincing causal inference methods. Yet, determinants for these exaggeration and power issues remain understudied. Focusing on causal methods, I argue in this paper that exaggeration may be exacerbated by the foundational component of causal inference: the fact that it only leverages subsets of the variation. Although it enables to avoid confounding, it also reduces precision thus increasing the risk of exaggeration. The same aspect that makes causal identification strategies credible can create another type of bias. A systematic reporting of statistical power calculations and analysis of the variation actually used for identification could help gauge the risk of falling into this exaggeration trap.

References

Abadie, A. (2020), 'Statistical Nonsignificance in Empirical Economics', American Economic Review: Insights **2**(2), 193–208. 2, 11

Amrhein, V., Trafimow, D. and Greenland, S. (2019), 'Inferential Statistics as Descriptive Statistics: There Is No Replication Crisis if We Don't Expect Replication', The American Statistician **73**(sup1), 262–270. 10

Andrews, I. and Kasy, M. (2019), 'Identification of and Correction for Publication Bias', American Economic Review **109**(8), 2766–2794. 2, 11

Angrist, J. D. and Pischke, J.-S. (2009), Mostly Harmless Econometrics: An Empiricist's Companion, 1 edition edn, Princeton University Press, Princeton. 7

Angrist, J. D. and Pischke, J.-S. (2014), Mastering 'Metrics: The Path from Cause to Effect, Princeton University Press. 7

Aronow, P. M. and Samii, C. (2016), 'Does Regression Produce Representative Estimates of Causal Effects?', American Journal of Political Science **60**(1), 250–267. 9

Athey, S. and Imbens, G. (2016), 'The Econometrics of Randomized Experiments', arXiv:1607.00698 [econ, stat] . 7

Bagilet, V. (2023), 'Accurately Estimating Relatively Small Effects: Air Pollution and Health'. 8

Black, B. S., Hollingsworth, A., Nunes, L. and Simon, K. I. (2021), Simulated Power Analyses for Observational Studies: An Application to the Affordable Care Act Medicaid Expansion, SSRN Scholarly Paper ID 3368187, Social Science Research Network, Rochester, NY. 5, 6, 7

Brodeur, A., Cook, N. and Heyes, A. (2020), 'Methods Matter: P-Hacking and Publication Bias in Causal Analysis in Economics', American Economic Review **110**(11), 3634–3660. 2, 6, 11

Button, K. S., Ioannidis, J. P. A., Mokrysz, C., Nosek, B. A., Flint, J., Robinson, E. S. J. and Munafò, M. R. (2013), 'Power failure: Why small sample size undermines the reliability of neuroscience', Nature Reviews Neuroscience **14**(5), 365–376. 2

Camerer, C. F., Dreber, A., Forsell, E., Ho, T.-H., Huber, J., Johannesson, M., Kirchler, M., Almenberg, J., Altmejd, A., Chan, T., Heikensten, E., Holzmeister, F., Imai, T., Isaksson, S., Nave, G., Pfeiffer, T., Razen, M. and Wu, H. (2016), 'Evaluating replicability of laboratory experiments in economics', Science **351**(6280), 1433–1436. 2, 6, 11

Card, D. (1993), Using Geographic Variation in College Proximity to Estimate the Return to Schooling, Working Paper 4483, National Bureau of Economic Research. 8

Chang, A. C. and Li, P. (2022), 'Is Economics Research Replicable? Sixty Published Papers From Thirteen Journals Say "Often Not"', Critical Finance Review **11**. 2

Christensen, G. and Miguel, E. (2018), 'Transparency, Reproducibility, and the Credibility of Economics Research', Journal of Economic Literature **56**(3), 920–980. 6, 10

Cinelli, C. and Hazlett, C. (2020), 'Making sense of sensitivity: Extending omitted variable bias', Journal of the Royal Statistical Society: Series B (Statistical Methodology) **82**(1), 39–67. 5, 6

Cunningham, S. (2021), Causal Inference: The Mixtape, Yale University Press. 7

Deaton, A. and Cartwright, N. (2018), 'Understanding and misunderstanding randomized controlled trials', Social Science & Medicine **210**, 2–21. 5

Duflo, E., Glennerster, R. and Kremer, M. (2007), Using Randomization in Development Economics Research: A Toolkit, in T. P. Schultz and J. A. Strauss, eds, 'Handbook of Development Economics', Vol. 4, Elsevier, pp. 3895–3962. 7

Gelman, A. (2020), Regression and Other Stories, Cambridge University Press, Cambridge New York, NY Port Melbourne, VIC New Delhi Singapore. 5, 7

Gelman, A. and Carlin, J. (2014), 'Beyond Power Calculations: Assessing Type S (Sign) and Type M (Magnitude) Errors', Perspectives on Psychological Science **9**(6), 641–651. 2, 5, 8, 11

Gelman, A. and Loken, E. (2013), 'The garden of forking paths: Why multiple comparisons can be a problem, even when there is no "fishing expedition" or "p-hacking" and the research hypothesis was posited ahead of time'. 10

Hernán, M. A. (2021), 'Causal analyses of existing databases: No power calculations required', Journal of Clinical Epidemiology p. S0895435621002730. 10

Hernán, M. A. and Robins, J. M. (2020), Causal Inference: What If, boca raton: chapman & hall/crc edn. 5

Hill, J. L. (2011), 'Bayesian Nonparametric Modeling for Causal Inference', Journal of Computational and Graphical Statistics **20**(1), 217–240. 7

Huntington-Klein, N. (2021), The Effect: An Introduction to Research Design and Causality, first edn, Chapman and Hall/CRC, Boca Raton. 7

Imbens, G. and Kalyanaraman, K. (2012), 'Optimal Bandwidth Choice for the Regression Discontinuity Estimator', The Review of Economic Studies **79**(3), 933–959. 5

Imbens, G. W. and Rubin, D. B. (2015), Causal Inference for Statistics, Social, and Biomedical Sciences: An Introduction, Cambridge University Press, Cambridge. 7

Ioannidis, J. P. A. (2008), 'Why Most Discovered True Associations Are Inflated', Epidemiology **19**(5), 640–648. 2, 11

Ioannidis, J. P. A., Stanley, T. D. and Doucouliagos, H. (2017), 'The Power of Bias in Economics Research', The Economic Journal **127**(605), F236–F265. 2, 6, 8, 11

Kasy, M. (2021), 'Of Forking Paths and Tied Hands: Selective Publication of Findings, and What Economists Should Do about It', Journal of Economic Perspectives **35**(3), 175–192. 6, 10

Linden, A. (2019), 'RETRODESIGN: Stata module to compute type-S (Sign) and type-M (Magnitude) errors', Boston College Department of Economics. 8

Lu, J., Qiu, Y. and Deng, A. (2019), 'A note on Type S/M errors in hypothesis testing', British Journal of Mathematical and Statistical Psychology **72**(1), 1–17. 2, 11

McConnell, B. and Vera-Hernandez, M. (2015), Going beyond simple sample size calculations: A practitioner's guide, Technical report, Institute for Fiscal Studies. 7

McShane, B. B., Gal, D., Gelman, A., Robert, C. and Tackett, J. L. (2019), 'Abandon Statistical Significance', The American Statistician **73**(sup1), 235–245. 10

Open Science Collaboration (2015), 'Estimating the reproducibility of psychological science', Science **349**(6251), aac4716. 2

Oster, E. (2019), 'Unobservable Selection and Coefficient Stability: Theory and Evidence', Journal of Business & Economic Statistics **37**(2), 187–204. 6

Ravallion, M. (2020), Should the Randomistas (Continue to) Rule?, Working Paper 27554, National Bureau of Economic Research. 5

Romer, D. (2020), 'In Praise of Confidence Intervals', AEA Papers and Proceedings **110**, 55–60. 10

Rosenbaum, P. R. (2002), Observational Studies, Springer Series in Statistics, Springer New York, New York, NY. 6

Rosenthal, R. (1979), 'The file drawer problem and tolerance for null results', Psychological Bulletin **86**(3), 638–641. 2, 11

Schell, T. L., Griffin, B. A. and Morral, A. R. (2018), Evaluating Methods to Estimate the Effect of State Laws on Firearm Deaths: A Simulation Study, Technical report, RAND Corporation. 6

Shadish, W. R., Cook, T. D. and Campbell, D. T. (2002), Experimental and Quasi-experimental Designs for Generalized Causal Inference, Houghton Mifflin. 7, 10

Simmons, J. P., Nelson, L. D. and Simonsohn, U. (2011), 'False-Positive Psychology: Undisclosed Flexibility in Data Collection and Analysis Allows Presenting Anything as Significant', Psychological Science **22**(11), 1359–1366. 10

Stommes, D., Aronow, P. M. and Sävje, F. (2021), 'On the reliability of published findings using the regression discontinuity design in political science', arXiv:2109.14526 [stat] . 5, 6, 8

Timm, A., Gelman, A. and Carlin, J. (2019), 'Retrodesign: Tools for Type S (Sign) and Type M (Magnitude) Errors'. 8

Young, A. (2021), 'Leverage, Heteroskedasticity and Instrumental Variables in Practical Application', p. 43. 6

Zwet, E. and Gelman, A. (2021), 'A Proposal for Informative Default Priors Scaled by the Standard Error of Estimates', The American Statistician pp. 1–9. 8

Zwet, E., Schwab, S. and Senn, S. (2021), 'The statistical properties of RCTs and a proposal for shrinkage', Statistics in Medicine **40**(27), 6107–6117. 8

Zwet, E. W. and Cator, E. A. (2021), 'The significance filter, the winner's curse and the need to shrink', Statistica Neerlandica **75**(4), 437–452. 2, 8, 11